

PETE GHIORSE

AI/ML PRODUCT LEADER | GROUP PRODUCT MANAGER | FOUNDER

New York City | Open to regular Bay Area travel | pmghiorse@gmail.com | peterghiorse.com | LinkedIn | GitHub

PROFILE

I lead AI/ML products, build production agents, and publish rigorous evaluations of how they behave. My work connects product strategy, cross-functional execution, hands-on prototyping, model evaluation, and the operating judgment required to move from an impressive demo to a trustworthy product.

EXPERIENCE

Capital One | Group Product Manager, AI/ML

2022 - Present

Lead product direction for ML-powered spend controls and fraud detection.

- Own product strategy and roadmap decisions for AI/ML-dependent experiences.
- Lead cross-functional work across engineering, data science, design, and risk partners.
- Translate model uncertainty and operating constraints into product boundaries and launch criteria.

Honeydew | Founder

2025 - Present

Build a multimodal family-coordination agent and use it as a production lab for model behavior, evaluation, and trust.

- Own discovery, product strategy, implementation, instrumentation, evaluation, and distribution.
- Turn voice, text, and photo input into calendars, lists, reminders, and shared household plans.
- Designed a 227-scenario restraint study across six models and 4,086 calls, with methods, failures, and limitations published.
- Built a separate 2,800-call model-selection benchmark and preserved it as a dated, prompt-coupled case study.

GiveTide | Co-Founder

2017 - 2022

Co-founded and ran a charitable-giving fintech product for five years before its acquisition in 2022.

- Led the product from concept through launch, operation, and acquisition.
- Worked across product, customer discovery, full-stack implementation, and company building.

EDUCATION

University of Richmond, Robins School of Business | B.S. Finance & Economics | Magna Cum Laude

PETE GHIORSE

AI/ML PRODUCT LEADER | GROUP PRODUCT MANAGER | FOUNDER

New York City | Open to regular Bay Area travel | pmghiorse@gmail.com | peterghiorse.com | LinkedIn | GitHub

SELECTED PUBLIC EVIDENCE

Agent restraint evaluation

June 24, 2026 | 227 scenarios | 6 models | 3 trials each | 4,086 calls | 14 parse failures disclosed

Measured act/ask/confirm/chat behavior on synthetic household scenarios. The useful result was not one leaderboard: safety, messy-input robustness, steerability, and judge disagreement separated the models.

github.com/pghio/agent-restraint-evals

Production model-selection benchmark

April 15, 2026 | 8 models | 35 scenarios | 10 trials each | 2,800 calls | temperature 0.2

Compared tool use, restraint, parsing, latency, and cost on one production-shaped workload. Public materials preserve the prompt-coupling caveat and exclude the proprietary frozen prompt and raw provider responses.

github.com/pghio/llm-agent-benchmark

LLM discoverability field note

90-day window | 13 GA4 sessions | 29 separate custom events | 10-query citation panel

Re-audited an overconfident interpretation, kept sessions and events as different units, and published the corrected query panel, aggregate observations, and denominator check without visitor-level data.

github.com/pghio/llm-discoverability-field-note

CORE CAPABILITIES

AI product leadership: strategy, roadmaps, decision boundaries, governance, and cross-functional execution for products that depend on uncertain model behavior.

Agent evaluation and safety: act/ask/confirm policy, destructive-action guard cases, messy-input robustness, and production-shaped evaluation design.

Multimodal systems: voice, text, images, tool orchestration, and the product decisions that make an agent useful outside a demo.

Evidence-backed communication: methods, limitations, negative results, corrections, and visual explanation that let readers inspect the reasoning.

WORKING PRACTICE

Build -> Instrument -> Evaluate -> Publish -> Revise. AI assists with research, red-teaming, implementation, and editing. I own the thesis, source selection, evaluation design, factual claims, accepted code, and final editorial judgment.

TARGET ROLES

Group Product Manager | Director of Product Management | Principal Product Manager | Head of AI Product